Exercise Briefing: Beam Search

Ki Hyun Kim

nlp.with.deep.learning@gmail.com

Project Implementation

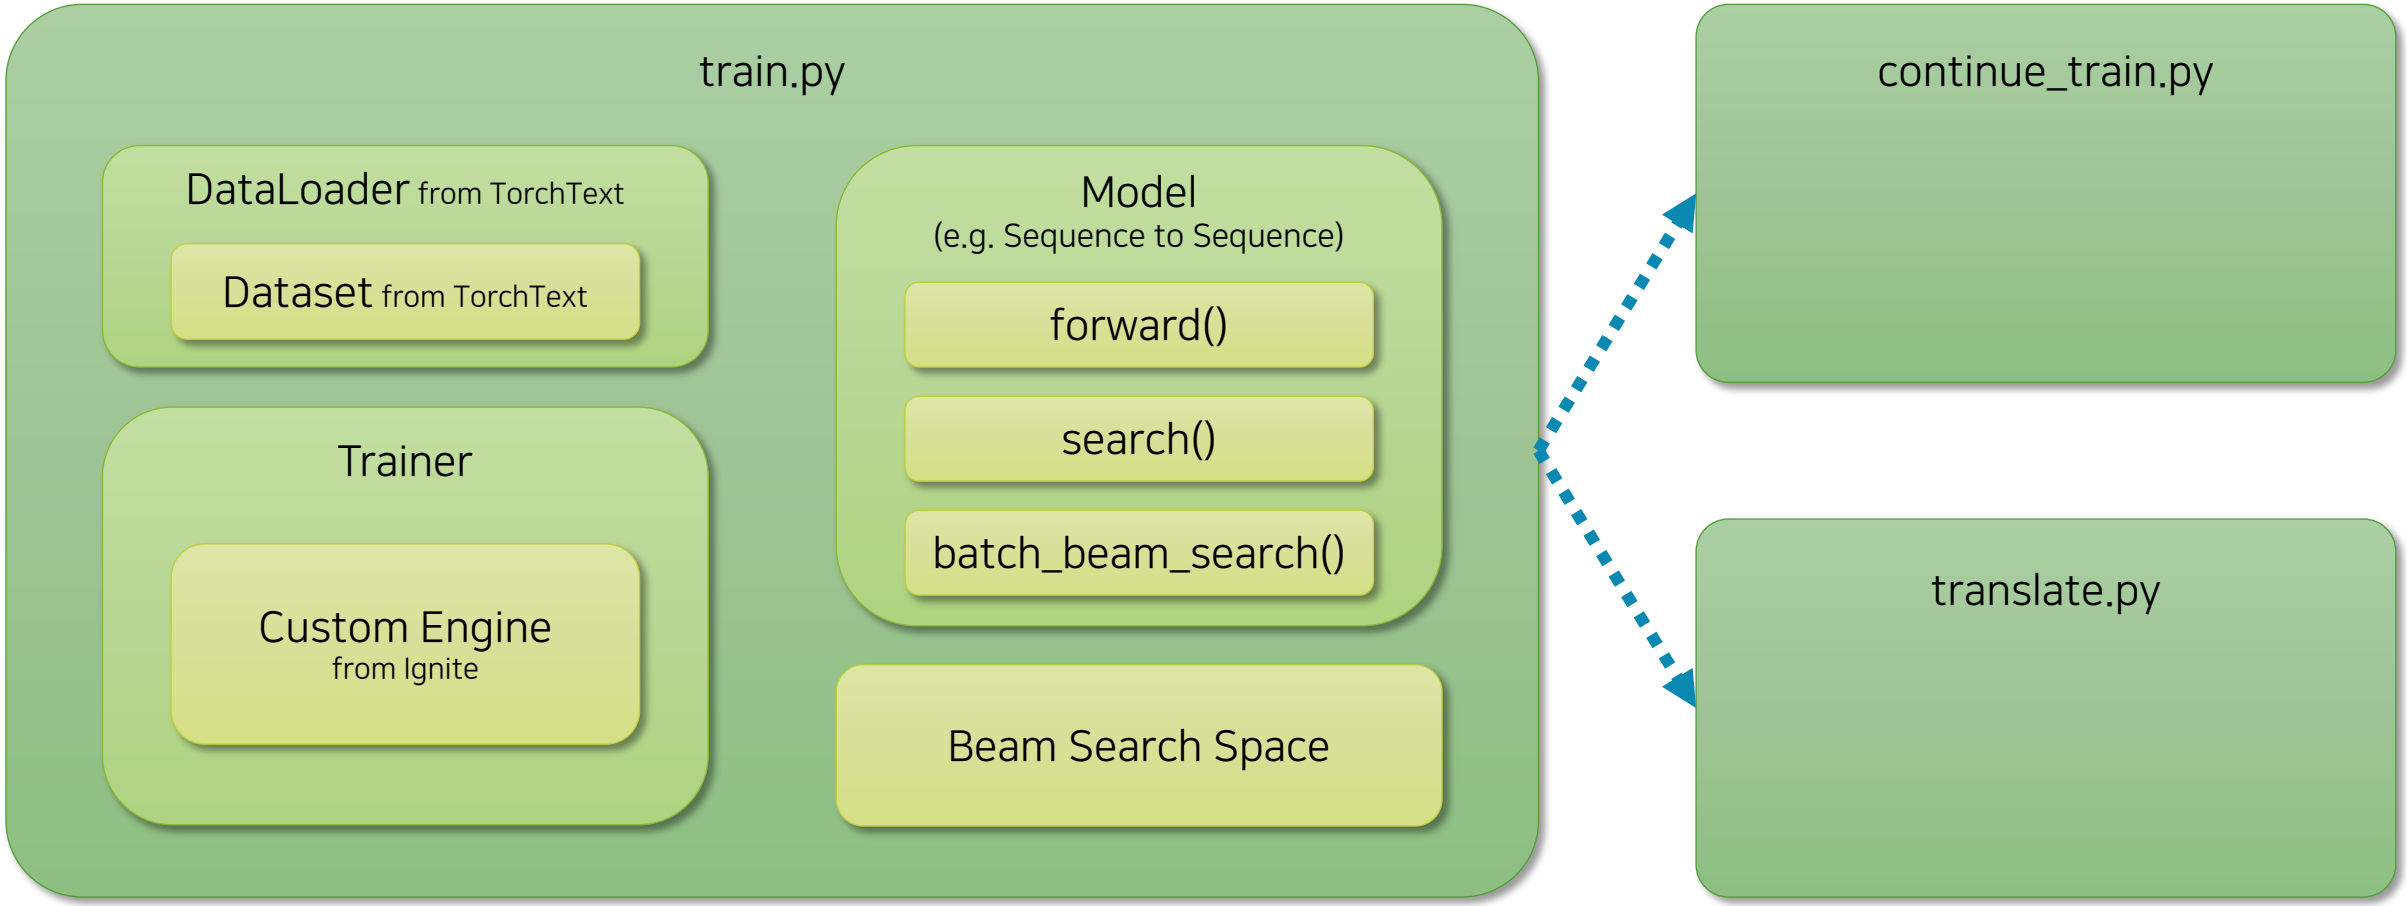

Mini-batch Parallelized Beam Search

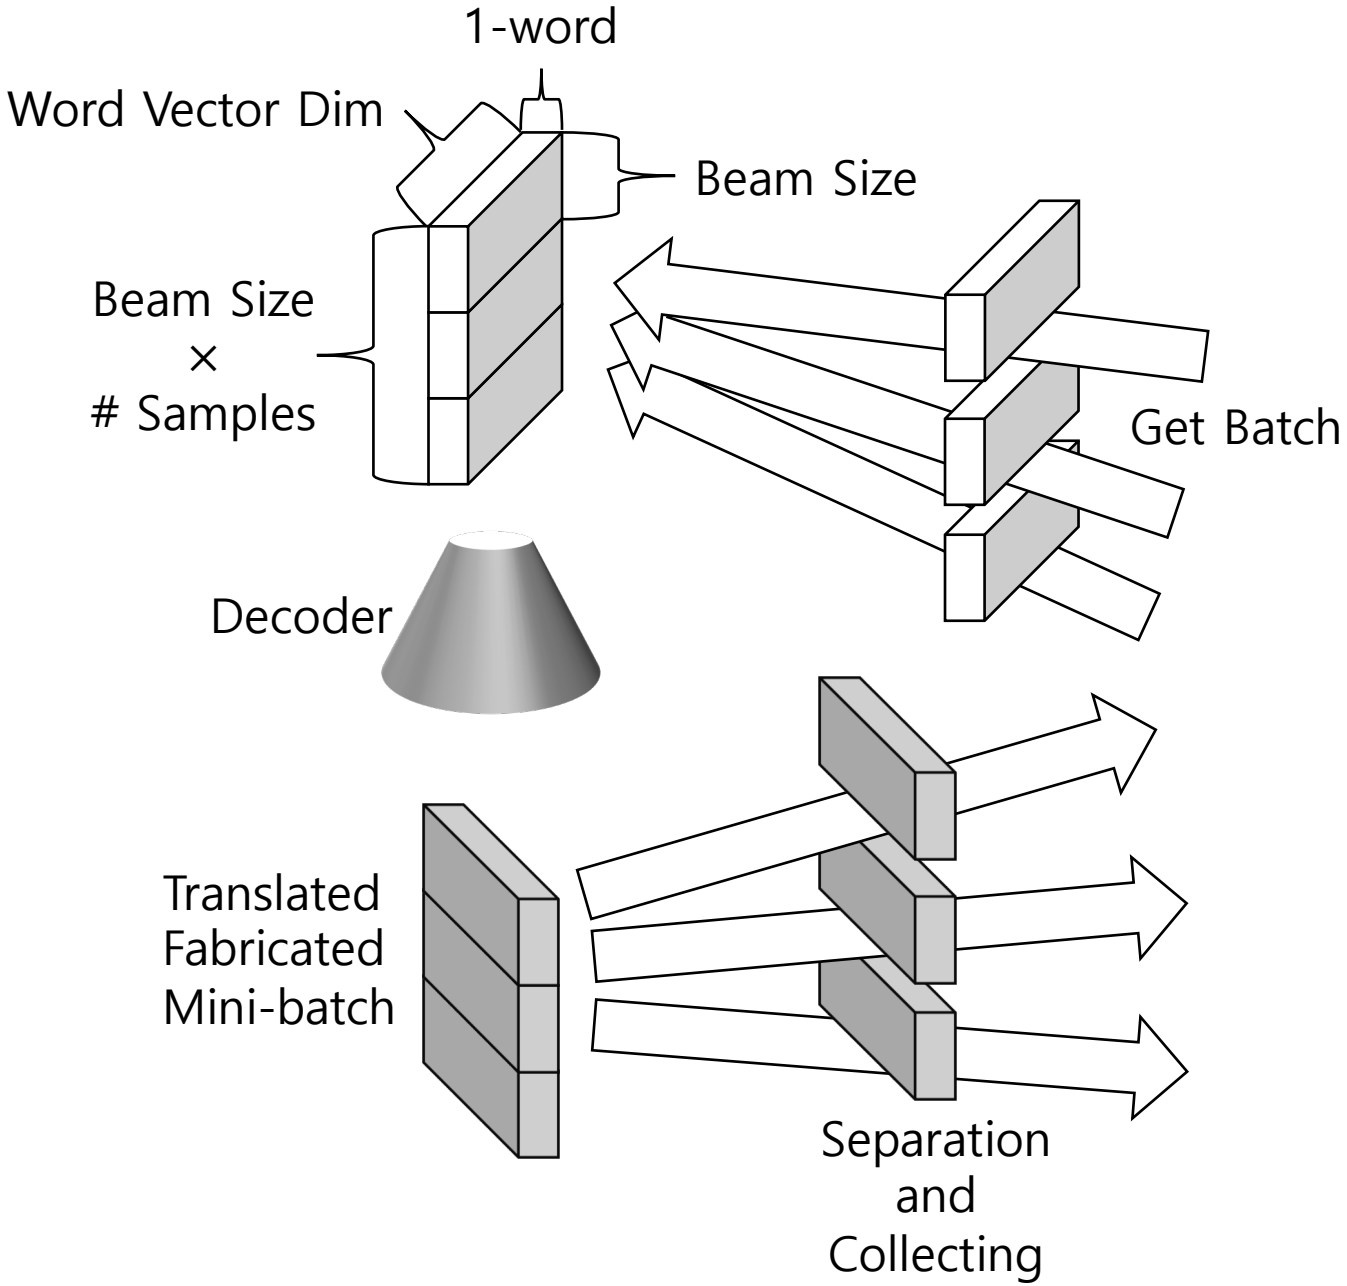

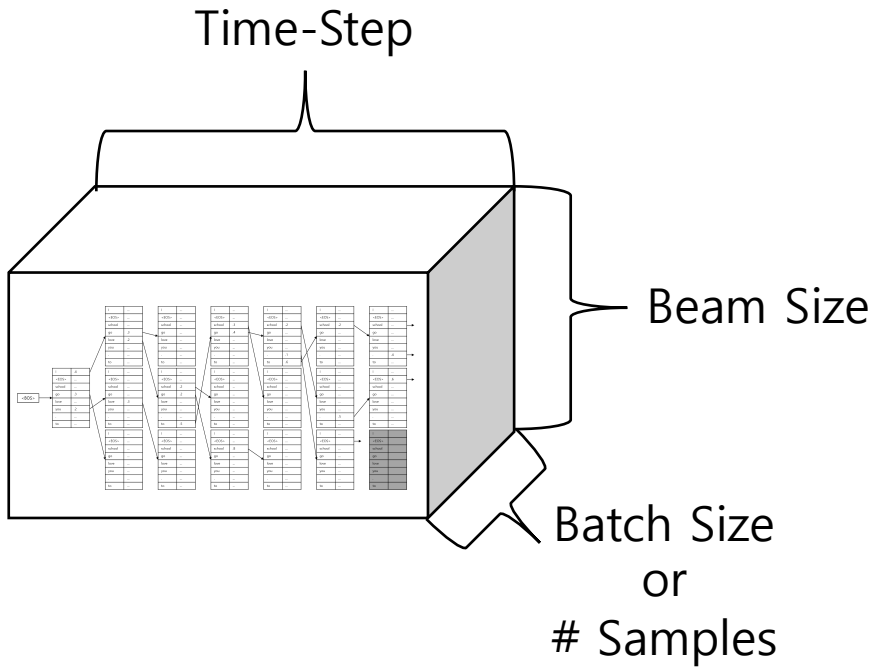

Mini-batch Parallelized Beam Search

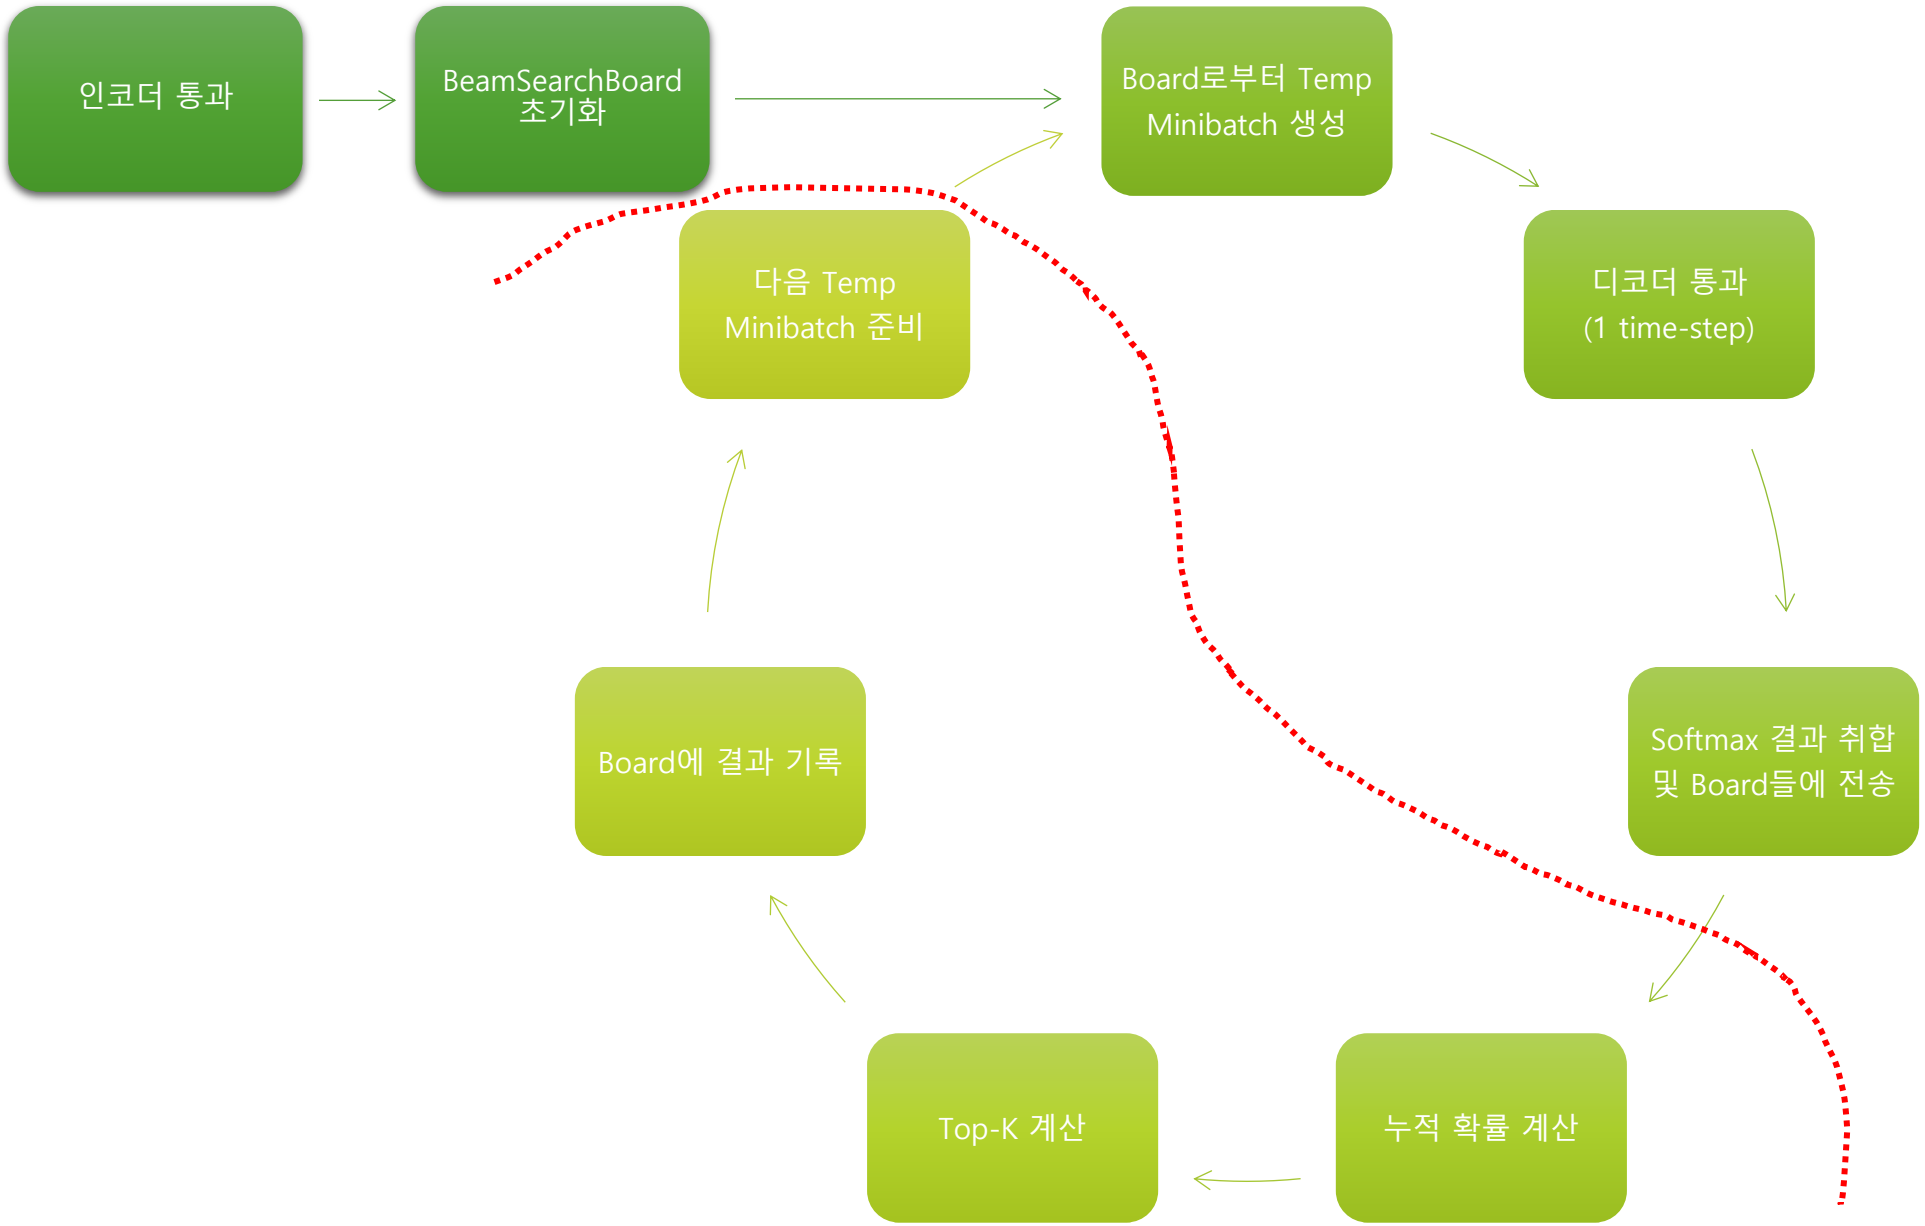

What we need to remember:

- From Board:
  - Input: x_t
  - Last hidden state: h_{t-1}
  - Last cell statec_{t-1}
  - Last H-tilde: $\tilde{h}_{t-1}$

- Expand to fake batch_size:
  - Entire output from encoder: $h_{1:n}^{enc}$
  - mask

- 주의: hidden & cell state와 H-tilde의 tensor shape 구성이 다름!