

# Exercise Briefing

Ki Hyun Kim

[nlp.with.deep.learning@gmail.com](mailto:nlp.with.deep.learning@gmail.com)

# Objective:

- 샘플링 및 reward & loss 계산하기
  - Search 함수 활용
- 기존의 모델(e.g. Seq2seq, Transformer)들을 그대로 활용할 수 있어야 함
  - 기왕이면 기존 Trainer도 최대한 재활용 하는 방향으로

# Review: REINFORCE with Baseline

Given policy  $\pi_{\theta}(a|s)$  and value function  $v_{\phi}(s)$ ,

For each episode:

Generate an episode  $s_0, a_0, r_1, \dots, s_{T-1}, a_{T-1}, r_T$  from  $\pi_{\theta}$ .

Loop: update  $\theta$  and  $\phi$  for each step of the episode  $t = 0, 1, \dots, T - 1$ :

$$G \leftarrow \sum_{k=t+1}^T \gamma^{k-t-1} r_k$$

$$\delta \leftarrow G - v_{\phi}(s_t)$$

$$\phi \leftarrow \phi + \eta^{\phi} \gamma^t \delta \nabla_{\phi} v_{\phi}(s_t)$$

$$\theta \leftarrow \theta + \eta^{\theta} \gamma^t \delta \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

# Equations

- What we will implement:

$$\mathcal{D} = \{x^i, y^i\}_{i=1}^N$$

$$\nabla_{\theta} \mathcal{L}(\theta) = \nabla_{\theta} \left( \sum_{i=1}^N \log P(\hat{y}_0^i | x^i; \theta) \times - \left( \text{reward}(\hat{y}_0^i, y^i) - \frac{1}{K} \sum_{k=1}^K \text{reward}(\hat{y}_k^i, y^i) \right) \right),$$

where  $\hat{y}^i \sim P(\cdot | x^i; \theta)$ .

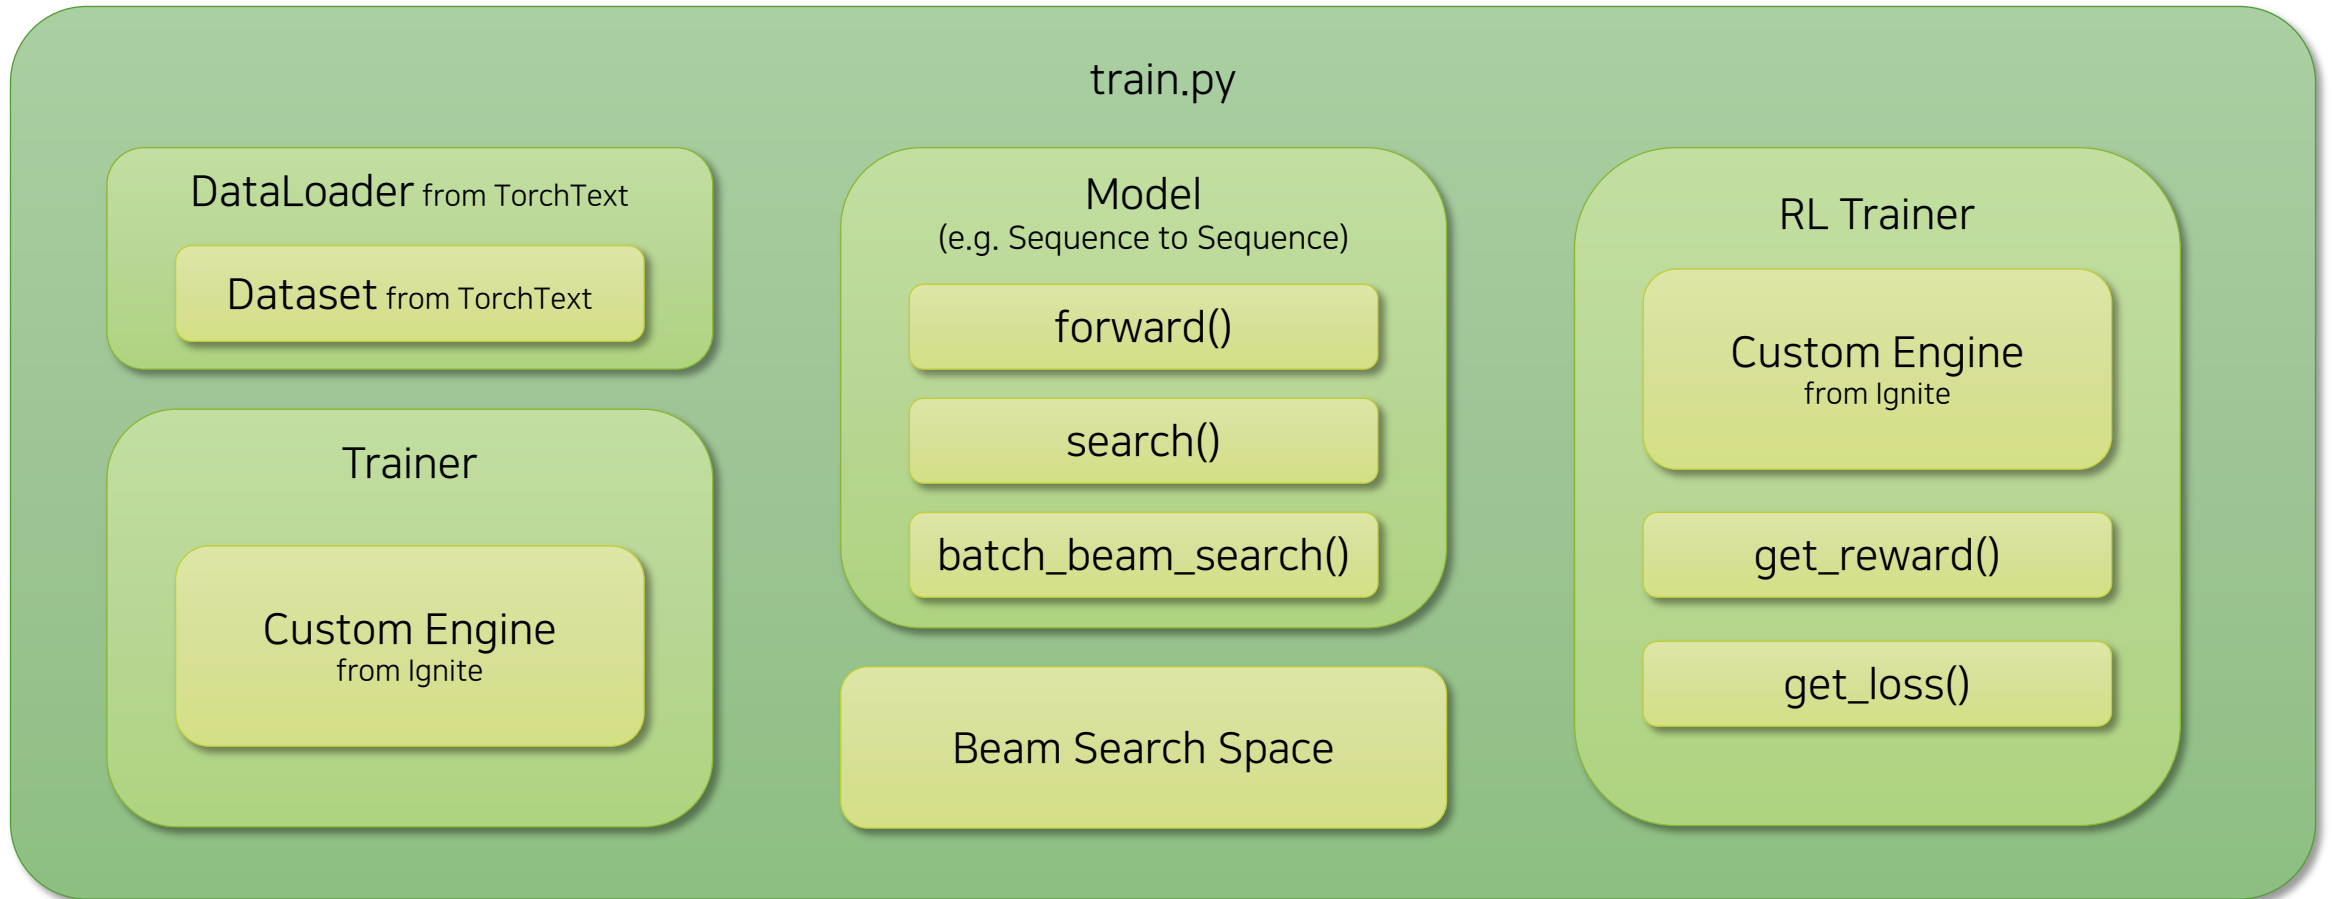
$$\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}(\theta)$$

# Equations

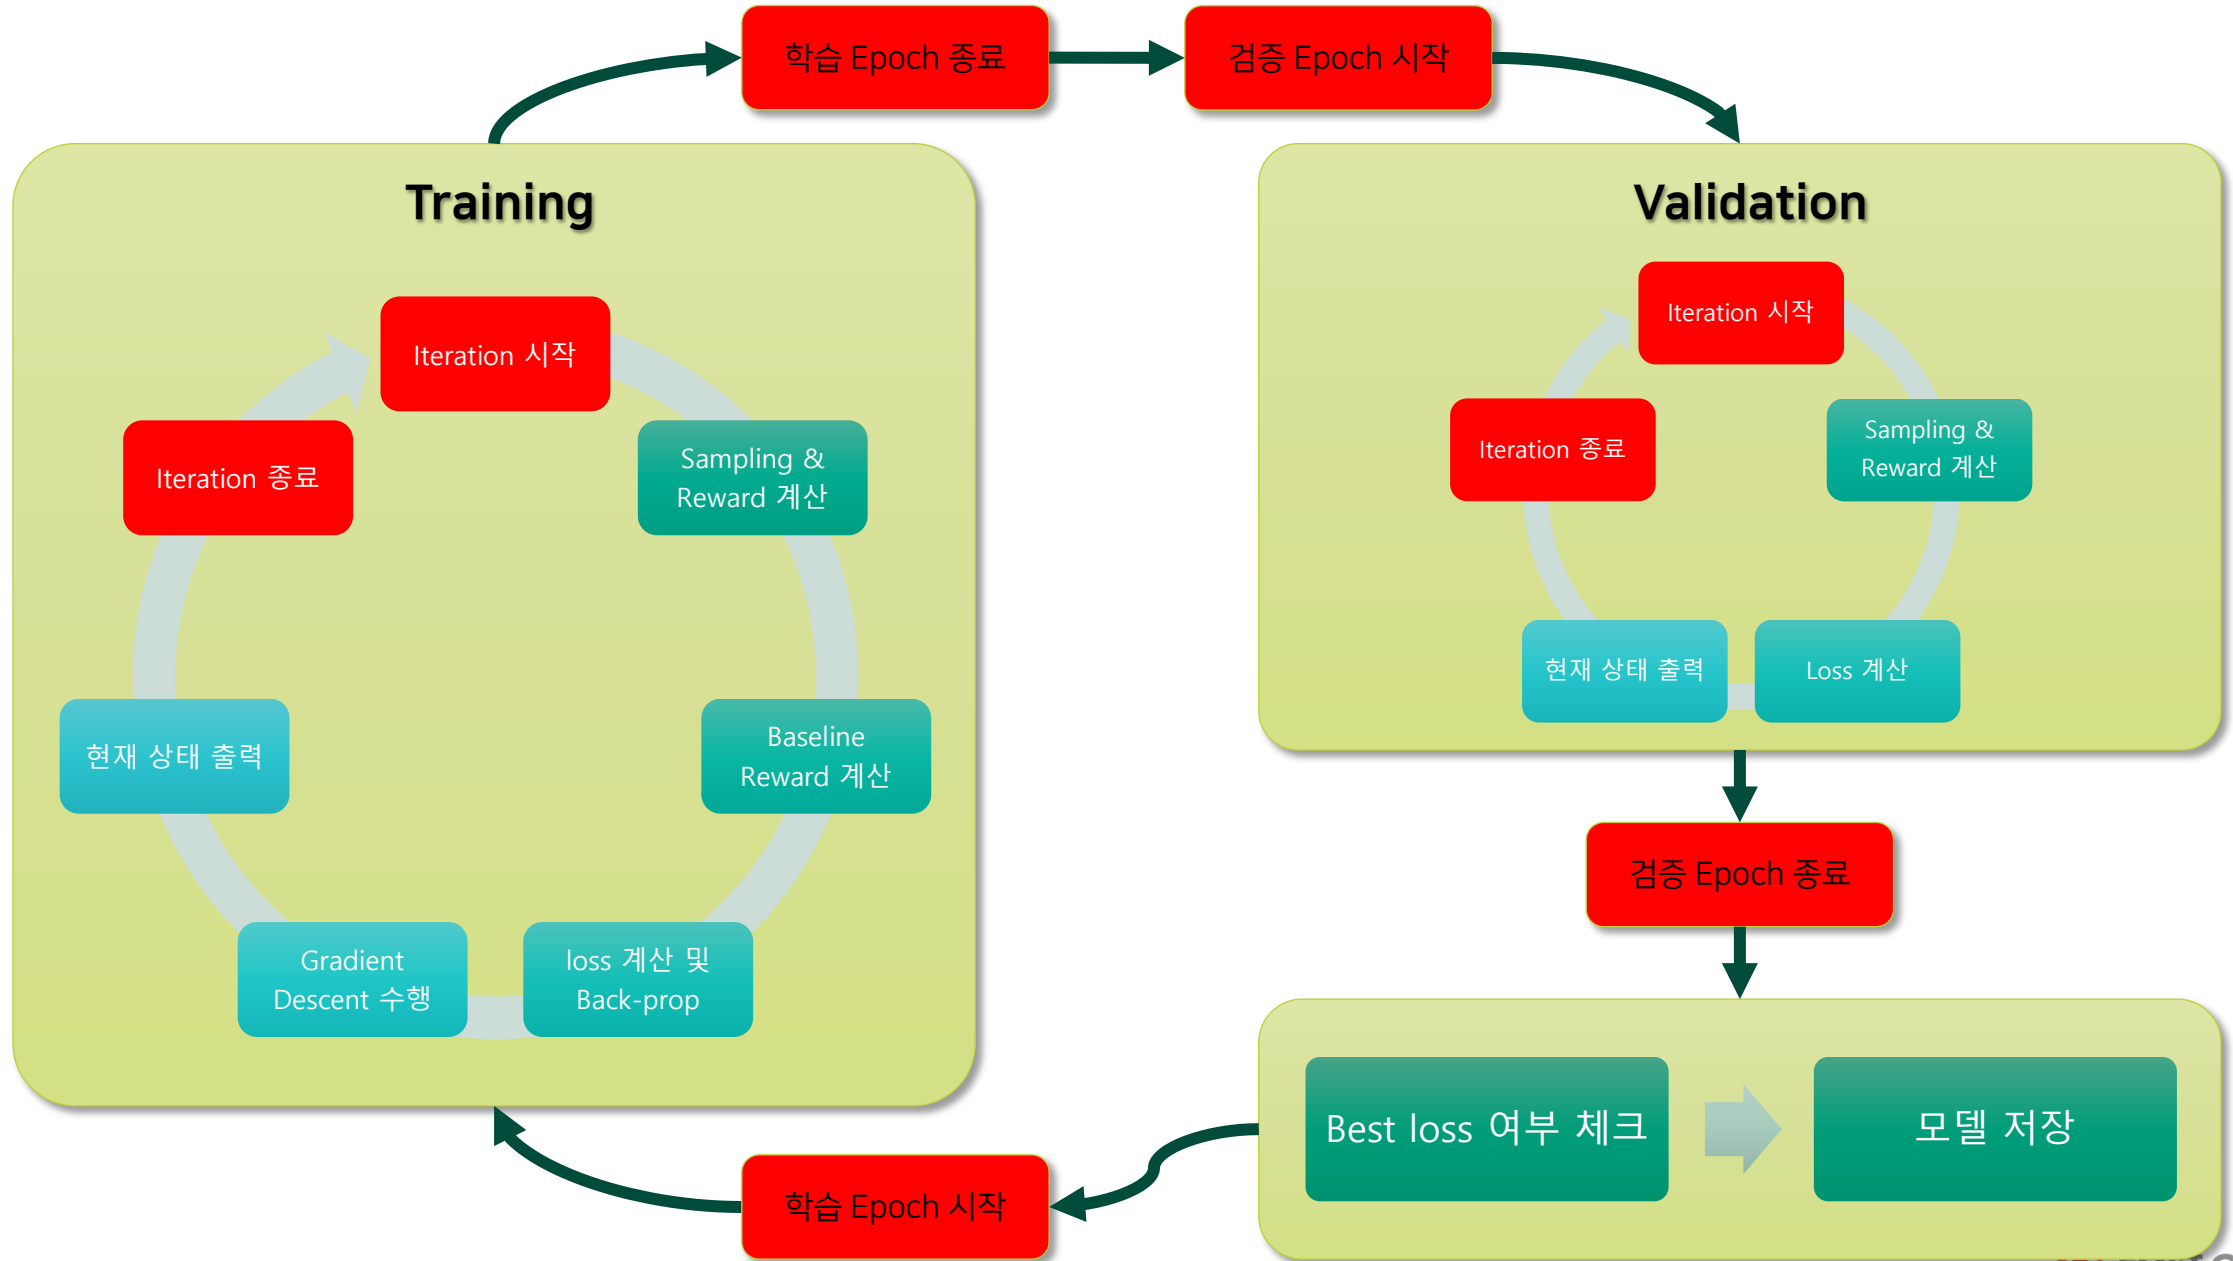
- How to implement:

$$\log P(\hat{y}|x; \theta) = \sum_{t=1}^{\tilde{m}} \hat{y}_t^\top \cdot \log f_\theta(x, \hat{y}_{<t})$$

# Project Implementation



# Training Procedure



# Weird & Practical Tip

- Sort() is much faster? than Topk() in CUDA
  - <https://github.com/pytorch/pytorch/issues/22812>

```
import time
import torch

data = torch.rand(2000000, dtype=torch.float32, device=torch.device('cuda:0'))
num_topk = 1000

def topk1():
    return data.topk(num_topk, sorted=False)

def topk2():
    sort, idx = data.sort(descending=True)
    return sort[:num_topk], idx[:num_topk]

def benchmark(f, iter, warmup):
    for k in range(warmup): f()
    start = time.perf_counter()
    for k in range(iter): f()
    torch.cuda.synchronize()
    return time.perf_counter() - start

print(benchmark(topk1, 100, 3))
print(benchmark(topk2, 100, 3))
print(benchmark(topk1, 100, 3))
print(benchmark(topk2, 100, 3))
```



```
2.557645708322525
0.3197173401713371
2.5575470123440027
0.31977611407637596
```